

HW 1 due this Friday, Jan 30 (10:00pm)

Midterm next Tuesday (Feb 3) 9:30am
in classroom

Class 5: *Regular Expressions*

University of Virginia
CS3120: DMT2

<https://weikailin.github.io/cs3120-toc>
Wei-Kai Lin



Regular Expressions

Midterm review

Plan

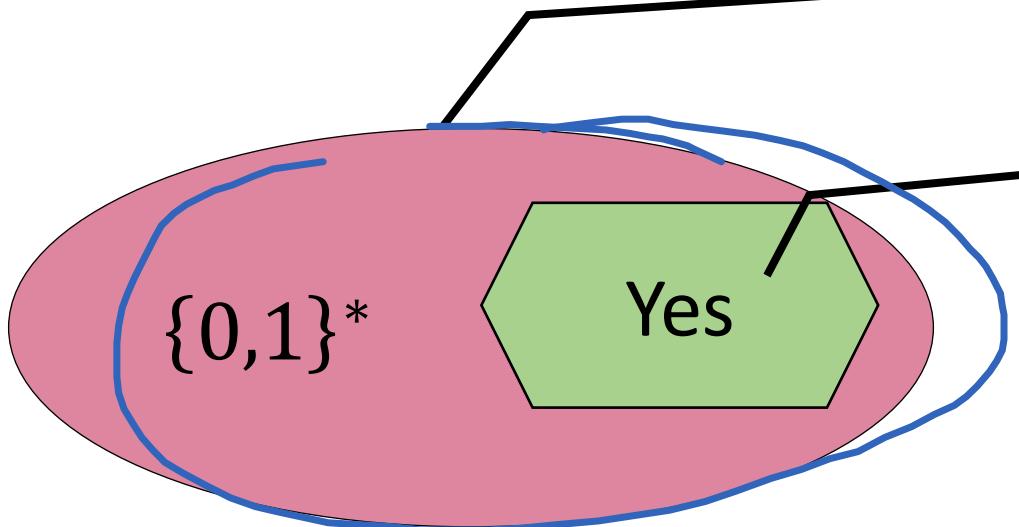
Today: Chapter 6 in the TCS book

https://introtcs.org/public/lec_05_infinite.html#regexpsec

Recap: Problems and Languages

Problem = Language = ~~Boolean n~~ Binary Function

Any binary string $x \in \{0,1\}^*$ is an *instance*.



The subset of Yes instances:
The *Problem*, or the *Language*,
or the Boolean function
that outputs 1

Class DFA

Recap: Deterministic Finite Automata

Definition 6.2 (Deterministic Finite Automaton)

A deterministic finite automaton (DFA) with C states over $\{0, 1\}$ is a pair (T, \mathcal{S}) with $T : [C] \times \{0, 1\} \rightarrow [C]$ and $\mathcal{S} \subseteq [C]$. The finite function T is known as the **transition function** of the DFA.

The set \mathcal{S} is known as the set of **accepting states**.

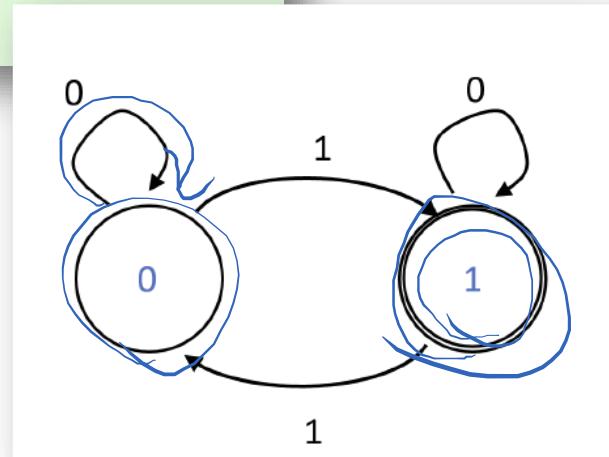
Let $F : \{0, 1\}^* \rightarrow \{0, 1\}$ be a Boolean function with the infinite domain $\{0, 1\}^*$. We say that (T, \mathcal{S}) **computes** a function $F : \{0, 1\}^* \rightarrow \{0, 1\}$ if for every $n \in \mathbb{N}$ and $x \in \{0, 1\}^n$, if we define $s_0 = 0$ and $s_{i+1} = T(s_i, x_i)$ for every $i \in [n]$, then

$$s_n \in \mathcal{S} \Leftrightarrow F(x) = 1$$

Goal: Same machine that computes on every binary string $x \in \{0, 1\}^*$.

Definition of DFA

Definition of DFA output



Regular Expressions

A (seemingly) completely different way of dealing with infinite languages
(i.e., functions on infinite input sets)

Motivation and example

Suppose we want to describe key-words for search

Simple examples: “book” or “February” or “a” or “10”

Intense example: a sorted sequence of integers

Unbound

Maybe crazy: a very large integer that is prime

Can we support all above?

Do we want to support all?

Motivation: simple rules

How about patterns with simple rules.

Example: we want a string of zeros only or ones only of length at least 2

00000
11111
We denote them as: $(00(0^*)|11(1^*))$

- 00 simply means string 00

- 0^* means repeating 0 zero, or one, or two, or ... times.

- | means OR

and or if

Regular Expression is Widely Used

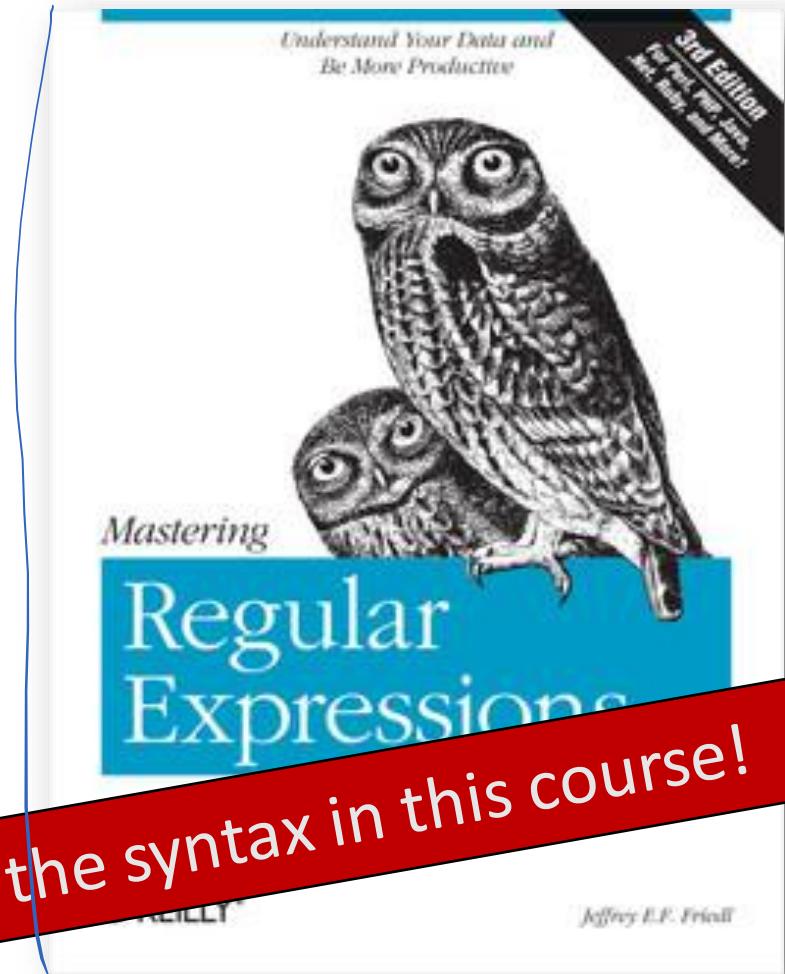
Examples:

*.pdf: any string ends with '.pdf'

?pdf: any string ends with 'pdf'
that is exactly 5 chars

Wildcards is a subset of regular
expressions.

This is Not the syntax in this course!



Writing “regular expressions” using “regular operations”

Definition 6.6 (Regular expression)

A *regular expression* e over an alphabet Σ is a string over $\Sigma \cup \{((), |, *, \emptyset, "")\}$ that has one of the following forms:

1. $e = \sigma$ where $\sigma \in \Sigma$
2. $e = (e'|e'')$ where e', e'' are regular expressions.
3. $e = (e')(e'')$ where e', e'' are regular expressions. (We often drop the parentheses when there is no danger of confusion and so write this as $e' e''$.)
4. $e = (e')^*$ where e' is a regular expression.

Often $\Sigma = \{0,1\}$ to be simple

Finally we also allow the following “edge cases”: $e = \emptyset$ and $e = ""$. These are the regular expressions corresponding to accepting no strings, and accepting only the empty string respectively.

(Informal) Intuition

- $\Sigma = \{0,1\}$: exact match
- \emptyset : matches nothing
- $""$: matches only empty string
- Special symbols: $($, $)$, $*$
- $(e') \mid (e'')$: e' matches OR e'' matches
- $(e')(e'')$: e' matches the prefix and e'' matches the suffix
- $(e')^*$: matches a string “repeatedly” for 0 or more times

concat

induct

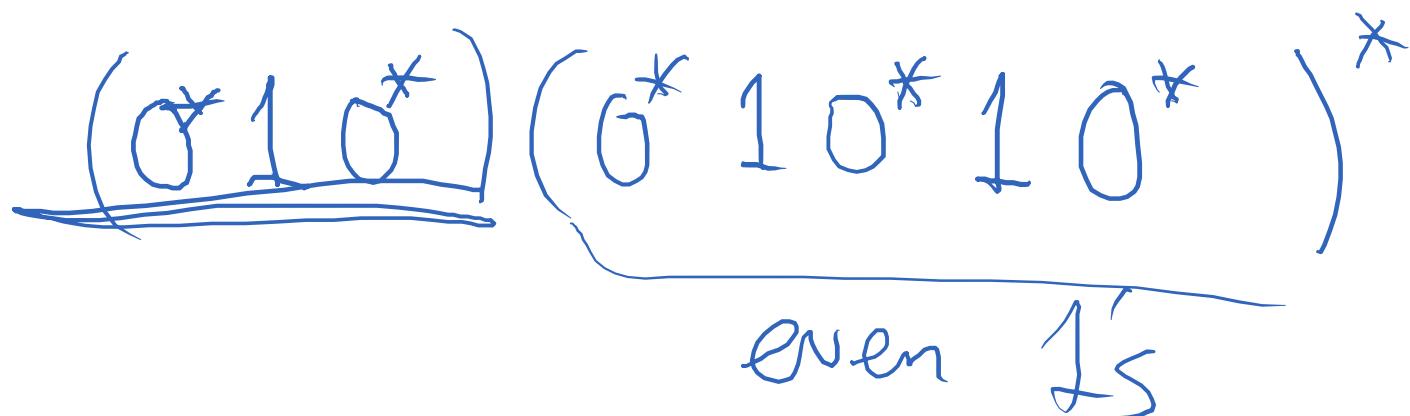
base

Example: XOR

For $x = x_1 \dots x_n$ let $\underline{\text{XOR}(x)} = \underline{x_1 \oplus x_2 \dots \oplus x_n}$

$$\text{XOR}(x) = 1 \text{ iff }$$

Observation: there are **odd** number of '1's in x



Parenthesis and Precedence

- Drop parentheses when inferred from context
- highest precedence to
*, then
concatenation, and then
OR
- $00^* \mid 11$ instead of $((0)(0^*)) \mid ((1)(1))$

Example: prefix and suffix

Alphabet $\Sigma = \{0,1\}$. All strings with:

Prefix: $|10|$

Suffix: $|000$

$$e = |10| (0|1)^* |000$$

$|10|000$

$(0|1)^*$

$(\cdot)^*$

N

Y

So far,
Reg. Exp. is just a string of
{alphabet and special symbols}

Evaluating Regular Expressions

For every regular expression e ,
there is a corresponding function $\Phi_e: \{0,1\}^* \rightarrow \{0,1\}$

Such that

$$\underbrace{\Phi_e(x)}_{\text{if } \underbrace{x}_{\text{matches}} \underbrace{e}} = 1$$



Defining Φ_e defines the evaluation of e .

Recursive definition, but cumbersome.

Definition 6.6 (Regular expression)

A regular expression e over an alphabet Σ is a string over $\Sigma \cup \{((), |, *, \emptyset, "")\}$ that has one of the following forms:

0. $e = \emptyset$, or $e = ""$

$$\Phi_{\emptyset}(x) = 0$$

$$\Phi_{""}(x) = \begin{cases} 1 & \text{iff } x = "" \\ 0 & \text{ow.} \end{cases}$$

1. $e = \sigma \in \{0,1\}$

$$\Phi_0, \Phi_1(x) = \begin{cases} 1 & \text{iff } x = "1" \\ 0 & \text{D.W.} \end{cases}$$

2. $e = (e'|e'')$

↑

$$\Phi_{e'}, \Phi_{e''}$$

$$\Phi_e(x) =$$

$$\Phi_{e'}(x) = 1 \text{ OR } \Phi_{e''}(x) = 1$$

e', e'' are regular expressions

Definition 6.6 (Regular expression)

A regular expression e over an alphabet Σ is a string over $\Sigma \cup \{((), |, *, \emptyset, "")\}$ that has one of the following forms:

3. $e = (e')(e'')$ $\Phi_e(x) = 1$ iff $\exists x_1, x_2$ s.t. $x_1 x_2 = x$
and $\Phi_{e'}(x_1) = 1$ and $\Phi_{e''}(x_2) = 1$
 $\vdash 1$

4. $e = (e')^*$ $\Phi_e(x) = 1$ if $\exists x_1, x_2, \dots, x_k$ s.t. $x_1 x_2 \dots x_k = x$

$\Phi_{e'}(x_i) = 1$ all $i = 1 \dots k$

$x_1 = " "$? ~~$k \in \mathbb{N}$~~

$|x| = n \in \mathbb{N}$
 ~~$k \leq n$~~ $k \in \mathbb{N}$

e', e'' are regular expressions

Definition 6.7 (Matching a regular expression)

Let e be a regular expression over the alphabet Σ . The function $\Phi_e : \Sigma^* \rightarrow \{0, 1\}$ is defined as follows:

1. If $e = \sigma$ then $\Phi_e(x) = 1$ iff $x = \sigma$.
2. If $e = (e'|e'')$ then $\Phi_e(x) = \Phi_{e'}(x) \vee \Phi_{e''}(x)$ where \vee is the OR operator.
3. If $e = (e')(e'')$ then $\Phi_e(x) = 1$ iff there is some $x', x'' \in \Sigma^*$ such that x is the concatenation of x' and x'' and $\Phi_{e'}(x') = \Phi_{e''}(x'') = 1$. $\xrightarrow{\text{N}}$ " " $\Phi_{e'}(" ") = 1$
4. If $e = (e')^*$ then $\Phi_e(x) = 1$ iff there is some $k \in \mathbb{N}$ and some $x_0, \dots, x_{k-1} \in \Sigma^*$ such that x is the concatenation $x_0 \cdots x_{k-1}$ and $\Phi_{e'}(x_i) = 1$ for every $i \in [k]$.
5. Finally, for the edge cases Φ_\emptyset is the constant zero function, and $\Phi_{\cdot\cdot\cdot}$ is the function that only outputs 1 on the empty string "".

We say that a regular expression e over Σ *matches* a string $x \in \Sigma^*$ if $\Phi_e(x) = 1$.

$\Phi_{e'}(" ") = 1$
↓
Don't consider

An (Python) algorithm evaluates regular expression

Python: 're'

10+ pages...

Regular Expression Syntax

A regular expression (or RE) specifies a set of strings that matches it; the functions in this module let you check if a particular string matches a given regular expression (or if a given regular expression matches a particular string, which comes down to the same thing).

The special characters are:

- .
- (Dot.) In the default mode, this matches any character except a newline. If the [DOTALL](#) flag has been specified, this matches any character including a newline. `(?s:.)` matches any character regardless of flags.
- ^
- (Caret.) Matches the start of the string, and in [MULTILINE](#) mode also matches immediately after each newline.
- \$
- Matches the end of the string or just before the newline at the end of the string, and in [MULTILINE](#) mode also matches before a newline. `foo` matches both 'foo' and 'foobar', while the regular expression `foo$` matches only 'foo'. More interestingly, searching for `foo.$` in '`foo1\nfoo2\n`' matches 'foo2' normally, but 'foo1' in [MULTILINE](#) mode; searching for a single \$ in '`foo\n`' will find two (empty) matches: one just before the newline, and one at the end of the string.
- *
- Causes the resulting RE to match 0 or more repetitions of the preceding RE, as many as possible. `ab*` will match 'a', 'ab', or 'a' followed by any number of 'b's.
- +
- Causes the resulting RE to match 1 or more repetitions of the preceding RE, as many as possible. `ab+` will match 'a' followed by any number of 'b's.

This is Not the syntax in this course!

Syntactic Sugar

Useful in practice,
but **not** used in cs3120

- Large alphabet:
Digits $0, 1, \dots, 9$. Letters a, b, \dots, z . Punctuations $\text{; } \text{. } \text{? } \text{(} \text{)}$

- Many special symbols:
Any char: $.$ Any digit: $\backslash d$
Any in list: $[abc\dots]$
Once or more: $(e')^+$
Constant repetition: $(e')\{n\}, (e')\{n,m\}$

$\neg (e')$

Negation? Stay tuned...

Complexity class: Regular Functions

Regular Functions / Language

Definition:

We call a Boolean function $F: \{0,1\}^* \rightarrow \{0,1\}$ is **regular**,
If $F = \Phi_e$ for some regular expression e .

Equivalently, a language $L \subseteq \{0,1\}^*$ is **regular**
if and only if there is a regular expression e such that
 $x \in L$ iff e matches x .

Complexity class: Regular Functions

Definition:

Let Reg-Fun be the set of all regular functions.

By definition:

For every $F \in \text{Reg-Fun}$, there exists a regular expression e such that $\Phi_e = F$.

Theorem: Reg-Fun = DFA-Comp

Theorem 6.17 (DFA and regular expression equivalency)

Let $F : \{0, 1\}^* \rightarrow \{0, 1\}$. Then F is regular if and only if there exists a DFA (T, S) that computes F .

reg exp iff DFA

Definitions:

Reg-Fun: the set $\{f \mid f = \Phi_e \text{ for some reg. exp. } e\}$

DFA-Comp: the set $\{f \mid f \text{ is computed by some DFA } M\}$

Theorem:

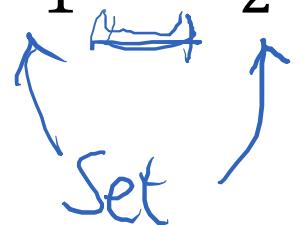
Reg-Fun = DFA-Comp (stay tuned for proof)

Theorem:

Reg-Fun = DFA-Comp

class \equiv set of func

For any two complexity classes C_1 and C_2 ,
what do we mean $C_1 = C_2$?



$A = B$ iff $A \subseteq B$
 $B \subseteq A$

What are C_1 and C_2 as math objects?

Set of funcs

Thm =

Interpret: $\text{Reg-Fun} \subseteq \text{DFA-Comp}$:

- For every $\underline{F} \in \text{Reg-Fun}$, $\underline{F} \in \text{DFA-Comp}$

$\exists \text{ DFA } M \text{ s.t. } M(x) = \underline{F}(x)$
all x

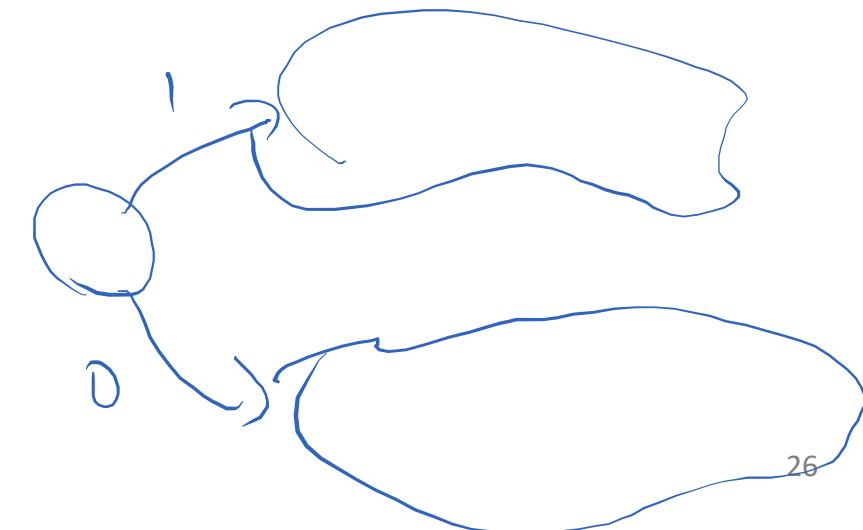
- Consequence: Every reg. exp. e , matching $\Phi_e(x)$ is computable in time $O(\underline{|x|})$ for all $x \in \{0,1\}^*$

why? bcs DFA M runs in $O(|x|)$ time

Interpret: Reg-Fun \supseteq DFA-Comp:

- For every $F \in \underline{\text{DFA-Comp}}$, $F \in \text{Reg-Fun}$
- Consequence: instead of writing DFA, enough to write a regular expression

$$e = (\underline{0}^*) \mid (\underline{1}^*)$$



Prepare for Midterm 0

Scope of Midterm 0

- Class 0 to Cantor's Thm in Class 4
- Quiz 0 to Quiz 2, HW 0 to HW 1
- Textbook Chapters 1 to 2 (some are not covered)

"~~prefix~~ prefix encoding"

Recap of Module 0

- Constructive definitions, natural numbers, binary strings, and induction
- Functions (bijection, surjection, ...)
- Sets, cardinality of sets, infinite, countable, uncountable
- Cantor's theorem

How to prepare?

- The HW and Quiz questions are closer
> Go through them
 - Slides are close in materials (but you may want
question)
 - Textbook exercises are good (some are not in scope)
 - Ask LLM for more related exercises
- format/material*

Find some exercise questions that are similar to the following.

Which one of these objects can be represented by a binary string?

An integer x

An undirected graph G

A directed graph H

All of the above.



That's a classic foundational question in computer science! It touches on the concept of **encoding**, which is the idea that almost any structured data can be flattened into a sequence of 0s and 1s for a computer to process.

Here are a few practice questions that explore similar themes—ranging from data structures to complexity and logic.

Format of Midterm 0

- We start 9:30am next Tuesday, Feb 3
- 20 minutes
- Closed book and everything, except for 1 page of letter-size double-sided sheet
- T/F, multiple choice, proof
- Simple calculation, reciting is not expected
(not even the sheet)

Purposes behind exam

- Grading: to meet the criteria of CS department
<https://uvacsadvising.org/policies.html#cs-department-grading-guidelines>
- WK's selfish goal:
 - to discover your potential in theoretical CS
 - (No need to be discouraged if it is too hard)

Logistics

- HW 2 will be posted in 24 hours about DFA and reg. exp., due next Friday
- We want to publish HW 1, but the earliest is next Monday
- WK will be in person on Feb 3.

